

# Experiment 2

Greenfield town is seeing rapid real estate development. A real estate company wants to predict house prices based solely on the size of the house (in sq ft) they have collected data from 2 recently sold houses.

Size (sqft)	Price (\$)
1200	300000
1500	320000
2000	450000
1800	400000
2500	550000
3000	600000
2200	480000
2700	700000

```
In [48]: import numpy as np
import matplotlib.pyplot as plt
# Data
sizes = np.array([1200, 1500, 2000, 1800, 2500, 3000, 2200, 2700])
prices = np.array([300000, 320000, 450000, 400000, 550000, 600000, 480000, 700000])
```

```
# Means
x_mean = sizes.mean()
y_mean = prices.mean()

# Compute slope m manually
numerator = np.sum((sizes - x_mean) * (prices - y_mean))
denominator = np.sum((sizes - x_mean)**2)
m = numerator / denominator

# Compute intercept b manually
b = y_mean - m * x_mean

# Predictions
y_pred = m * sizes + b

print(f"Slope (Coefficient): {m:.2f}")
print(f"Intercept: {b:.2f}")
```

Slope (Coefficient): 212.17  
 Intercept: 26789.65

```
In [49]: import numpy as np

# Data
sizes = np.array([1200, 1500, 2000, 1800, 2500, 3000, 2200, 2700])
prices = np.array([300000, 320000, 450000, 400000, 550000, 600000, 480000, 700000])
```

```
# Construct X matrix with a column of 1s
X = np.column_stack((np.ones(len(sizes)), sizes))

# y vector
y = prices.reshape(-1, 1)

# Closed-form solution  $\theta = (X^T X)^{-1} X^T y$ 
theta = np.linalg.inv(X.T @ X) @ X.T @ y

theta0 = theta[0, 0] # intercept
theta1 = theta[1, 0] # slope

print(f"Theta0 (Intercept): {theta0:.2f}")
print(f"Theta1 (Slope): {theta1:.2f}")
```

Theta0 (Intercept): 26789.65

Theta1 (Slope): 212.17

In [50]:


```
xpred = 1940
y_pred = m * xpred + b
print(f"Predicted price for size {xpred} sqft: {y_pred:.2f}")
```

Predicted price for size 1940 sqft: 438400.57

In [51]:

```
y_line = m * sizes + b

plt.scatter(sizes, prices, label='Data Points')
plt.plot(sizes, y_line, label='Regression Line')
plt.xlabel('Size (sqft)')
plt.ylabel('Price ($)')
plt.title('House Price Prediction')
plt.legend()
plt.show()
```



## Experiment 2.2

```
In [52]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LinearRegression

data=pd.read_csv('../assets/exp2_data.csv')
```

```
In [53]: data.head(5)
```

Out[53]:

	YearsExperience	Salary
<b>0</b>	1.1	39343.0
<b>1</b>	1.3	46205.0
<b>2</b>	1.5	37731.0
<b>3</b>	2.0	43525.0
<b>4</b>	2.2	39891.0

```
In [54]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# Prepare X and y
X = data['YearsExperience'].values.reshape(-1, 1)
y = data['Salary'].values

# Split into train/test sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# Train the model on training data
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on test data
y_pred = model.predict(X_test)

# Get slope and intercept
m = model.coef_[0]
b = model.intercept_

print(f"Slope (Coefficient): {m:.2f}")
print(f"Intercept: {b:.2f}")
```

Slope (Coefficient): 9423.82  
 Intercept: 25321.58

```
In [55]: y_line = model.predict(X) # prediction for the full dataset

plt.scatter(data['YearsExperience'], data['Salary'], color='blue', label='Data P
plt.plot(data['YearsExperience'], y_line, color='red', label='Regression Line')
```

```
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.title('Salary vs Experience')
plt.legend()
plt.show()
```



In [56]: `# Predict test values  
y_pred = model.predict(X_test)  
print(y_pred)`

```
[115790.21011287 71498.27809463 102596.86866063 75267.80422384  
55477.79204548 60189.69970699]
```

In [57]: `# Calculate R^2 score  
r2_score = model.score(X_test, y_test)  
print(f"R^2 Score: {r2_score:.4f}")`

```
R^2 Score: 0.9024
```

In [58]: `import pickle  
import os  
  
# Ensure the folder exists (optional but recommended)  
os.makedirs("../models", exist_ok=True)  
  
filename = '../models/experiment2-2.sav'  
  
with open(filename, 'wb') as f:  
 pickle.dump(model, f)`

In [59]: `loaded_model = pickle.load(open(filename, 'rb'))  
loaded_model`

Out[59]:

LinearRegression		
Parameters		
fit_intercept	True	
copy_X	True	
tol	1e-06	
n_jobs	None	
positive	False	

```
In [60]: result=loaded_model.predict([[15]])
print(result)
```

[166678.81285724]

```
In [61]: prediction_input=int(input("Enter years of experience: "))
print(loaded_model.predict([[prediction_input]]))
```

[336307.4886718]