

# Experiment 9: K-Means Clustering on Heart Disease Dataset

This notebook demonstrates K-Means clustering analysis on the Heart Disease dataset from UCI Machine Learning Repository.

## Objectives:

- Load and pre-process the Heart Disease dataset
- Determine optimal number of clusters using Silhouette score
- Reduce dimensionality to 2D using PCA
- Visualize clustered data and compare with actual labels

## Step 0: Import Required Libraries

Import necessary Python libraries for data manipulation, clustering, and visualization.

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.metrics import silhouette_score
import urllib.request
import warnings

warnings.filterwarnings("ignore")

# Set random seed for reproducibility
np.random.seed(42)
```

## Step 1: Load and Pre-process the Dataset

Load the Heart Disease dataset from UCI Repository. The dataset contains 14 attributes including:

- **Demographic:** age, sex
- **Symptoms:** cp (chest pain type), trestbps (resting blood pressure), chol (cholesterol)
- **Test results:** restecg, thalach, exang, oldpeak, slope, ca, thal
- **Target:** presence of heart disease (0-4, converted to binary)

```
In [3]: print("=" * 60)
print("STEP 1: Loading and Pre-processing Heart Disease Dataset")
print("=" * 60)

# Download the Heart Disease dataset from UCI
```

```

url = "https://archive.ics.uci.edu/ml/machine-learning-databases/heart-disease/p

try:
    # Download and Load the dataset
    column_names = [
        "age",
        "sex",
        "cp",
        "trestbps",
        "chol",
        "fbs",
        "restecg",
        "thalach",
        "exang",
        "oldpeak",
        "slope",
        "ca",
        "thal",
        "target",
    ]
    df = pd.read_csv(url, names=column_names, na_values="?")
    print("Dataset downloaded successfully from UCI Repository!")
except Exception as e:
    print(f"Error downloading dataset: {e}")
    print("Creating sample dataset for demonstration...")
    # Create synthetic data if download fails
    np.random.seed(42)
    n_samples = 303
    df = pd.DataFrame(
        {
            "age": np.random.randint(29, 77, n_samples),
            "sex": np.random.randint(0, 2, n_samples),
            "cp": np.random.randint(0, 4, n_samples),
            "trestbps": np.random.randint(94, 200, n_samples),
            "chol": np.random.randint(126, 564, n_samples),
            "fbs": np.random.randint(0, 2, n_samples),
            "restecg": np.random.randint(0, 3, n_samples),
            "thalach": np.random.randint(71, 202, n_samples),
            "exang": np.random.randint(0, 2, n_samples),
            "oldpeak": np.random.uniform(0, 6.2, n_samples),
            "slope": np.random.randint(0, 3, n_samples),
            "ca": np.random.randint(0, 4, n_samples),
            "thal": np.random.randint(0, 4, n_samples),
            "target": np.random.randint(0, 2, n_samples),
        }
    )

# Handle missing values
df = df.dropna()

print(f"Dataset shape: {df.shape}")
print("\nFirst 5 records:")
print(df.head())
print("\nDataset info:")
print(df.info())
print("\nTarget distribution:")
print(df["target"].value_counts())
print(f"\nMissing values: {df.isnull().sum().sum()}")

```

```
=====
STEP 1: Loading and Pre-processing Heart Disease Dataset
=====
Dataset downloaded successfully from UCI Repository!
Dataset shape: (297, 14)
```

First 5 records:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63.0	1.0	1.0	145.0	233.0	1.0	2.0	150.0	0.0	2.3	
1	67.0	1.0	4.0	160.0	286.0	0.0	2.0	108.0	1.0	1.5	
2	67.0	1.0	4.0	120.0	229.0	0.0	2.0	129.0	1.0	2.6	
3	37.0	1.0	3.0	130.0	250.0	0.0	0.0	187.0	0.0	3.5	
4	41.0	0.0	2.0	130.0	204.0	0.0	2.0	172.0	0.0	1.4	

	slope	ca	thal	target
0	3.0	0.0	6.0	0
1	2.0	3.0	3.0	2
2	2.0	2.0	7.0	1
3	3.0	0.0	3.0	0
4	1.0	0.0	3.0	0

Dataset info:

<class 'pandas.DataFrame'>

Index: 297 entries, 0 to 301

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	age	297 non-null	float64
1	sex	297 non-null	float64
2	cp	297 non-null	float64
3	trestbps	297 non-null	float64
4	chol	297 non-null	float64
5	fbs	297 non-null	float64
6	restecg	297 non-null	float64
7	thalach	297 non-null	float64
8	exang	297 non-null	float64
9	oldpeak	297 non-null	float64
10	slope	297 non-null	float64
11	ca	297 non-null	float64
12	thal	297 non-null	float64
13	target	297 non-null	int64

dtypes: float64(13), int64(1)

memory usage: 34.8 KB

None

Target distribution:

target

0	160
1	54
2	35
3	35
4	13

Name: count, dtype: int64

Missing values: 0

## Data Preparation

Separate features from target variable and standardize the features for clustering.

```
In [4]: # Separate features and target (convert target to binary: 0 = no disease, 1 = disease)
X = df.drop("target", axis=1)
y = (df["target"] > 0).astype(int) # Convert to binary classification

print(f"\nFeatures shape: {X.shape}")
print(f"Target distribution (binary):\n{y.value_counts()}")

# Standardize the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
print(f"\nFeatures standardized successfully!")
```

```
Features shape: (297, 13)
Target distribution (binary):
target
0    160
1    137
Name: count, dtype: int64

Features standardized successfully!
```

## Step 2: Apply K-Means and Determine Optimal k

Use the **Silhouette Score** to find the optimal number of clusters. The silhouette score ranges from -1 to 1, where:

- **1**: Samples are well-clustered and far from neighboring clusters
- **0**: Samples are on or very close to the decision boundary
- **-1**: Samples may have been assigned to the wrong cluster

```
In [5]: print("\n" + "=" * 60)
print("STEP 2: Finding Optimal Number of Clusters")
print("=" * 60)

# Test different values of k
k_range = range(2, 11)
silhouette_scores = []
inertias = []

for k in k_range:
    kmeans = KMeans(n_clusters=k, random_state=42, n_init=10)
    cluster_labels = kmeans.fit_predict(X_scaled)

    # Calculate silhouette score
    silhouette_avg = silhouette_score(X_scaled, cluster_labels)
    silhouette_scores.append(silhouette_avg)
    inertias.append(kmeans.inertia_)

print(
    f"k={k}: Silhouette Score = {silhouette_avg:.4f}, Inertia = {kmeans.inertia_}"
)
```

```

# Find optimal k
optimal_k = k_range[np.argmax(silhouette_scores)]
print(f"\nOptimal number of clusters (k) = {optimal_k}")
print(f"Best Silhouette Score = {max(silhouette_scores):.4f}")

# Apply K-Means with optimal k
kmeans_optimal = KMeans(n_clusters=optimal_k, random_state=42, n_init=10)
cluster_labels = kmeans_optimal.fit_predict(X_scaled)

print(f"\nCluster distribution:")
unique, counts = np.unique(cluster_labels, return_counts=True)
for cluster, count in zip(unique, counts):
    print(f" Cluster {cluster}: {count} samples")
=====

STEP 2: Finding Optimal Number of Clusters
=====
k=2: Silhouette Score = 0.1751, Inertia = 3193.45
k=3: Silhouette Score = 0.1298, Inertia = 2931.15
k=4: Silhouette Score = 0.1304, Inertia = 2743.77
k=5: Silhouette Score = 0.1127, Inertia = 2609.80
k=6: Silhouette Score = 0.1137, Inertia = 2507.28
k=7: Silhouette Score = 0.1013, Inertia = 2404.55
k=8: Silhouette Score = 0.1060, Inertia = 2318.45
k=9: Silhouette Score = 0.1000, Inertia = 2252.42
k=10: Silhouette Score = 0.1120, Inertia = 2206.82

```

Optimal number of clusters (k) = 2  
 Best Silhouette Score = 0.1751

Cluster distribution:  
 Cluster 0: 184 samples  
 Cluster 1: 113 samples

## Step 3: Dimensionality Reduction using PCA

Apply **Principal Component Analysis (PCA)** to reduce the 13-dimensional feature space to 2 dimensions for visualization. PCA transforms the data to a new coordinate system where the greatest variance lies on the first coordinate (PC1), the second greatest variance on the second coordinate (PC2), and so on.

```

In [6]: print("\n" + "=" * 60)
print("STEP 3: Dimensionality Reduction using PCA")
print("=" * 60)

# Apply PCA to reduce to 2 dimensions
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

print(f"Original dimensions: {X_scaled.shape[1]}")
print(f"Reduced dimensions: {X_pca.shape[1]}")
print(f"\nExplained Variance Ratio:")
print(
    f" PC1: {pca.explained_variance_ratio_[0]:.4f} ({pca.explained_variance_ratio_[
    ]})
    print(
        f" PC2: {pca.explained_variance_ratio_[1]:.4f} ({pca.explained_variance_ratio_[
    ]})"
)

```

```

print(
    f" Total: {sum(pca.explained_variance_ratio_):.4f} ({sum(pca.explained_vari
)
=====
STEP 3: Dimensionality Reduction using PCA
=====
Original dimensions: 13
Reduced dimensions: 2

Explained Variance Ratio:
PC1: 0.2370 (23.70%)
PC2: 0.1235 (12.35%)
Total: 0.3604 (36.04%)

```

## Step 4: Visualize the Clusters in 2D

Create comprehensive visualizations including:

1. **Elbow Method:** Shows the optimal k using Within-Cluster Sum of Squares (WCSS)
2. **Silhouette Scores:** Plot of silhouette scores for different k values
3. **K-Means Clusters:** Clusters visualized in PCA space with centroids
4. **Actual Labels:** Ground truth labels for comparison

```

In [7]: print("\n" + "=" * 60)
print("STEP 4: Visualizing Clusters")
print("=" * 60)

# Create a figure with multiple subplots
fig, axes = plt.subplots(2, 2, figsize=(15, 12))

# Plot 1: Elbow Method
axes[0, 0].plot(k_range, inertias, "bo-", linewidth=2, markersize=8)
axes[0, 0].set_xlabel("Number of Clusters (k)", fontsize=12)
axes[0, 0].set_ylabel("Inertia (WCSS)", fontsize=12)
axes[0, 0].set_title("Elbow Method", fontsize=14, fontweight="bold")
axes[0, 0].grid(True, alpha=0.3)
axes[0, 0].axvline(
    x=optimal_k, color="r", linestyle="--", label=f"Optimal k={optimal_k}"
)
axes[0, 0].legend()

# Plot 2: Silhouette Scores
axes[0, 1].plot(k_range, silhouette_scores, "go-", linewidth=2, markersize=8)
axes[0, 1].set_xlabel("Number of Clusters (k)", fontsize=12)
axes[0, 1].set_ylabel("Silhouette Score", fontsize=12)
axes[0, 1].set_title(
    "Silhouette Score vs Number of Clusters", fontsize=14, fontweight="bold"
)
axes[0, 1].grid(True, alpha=0.3)
axes[0, 1].axvline(
    x=optimal_k, color="r", linestyle="--", label=f"Optimal k={optimal_k}"
)
axes[0, 1].legend()

# Plot 3: K-Means Clusters in PCA Space
colors = [
    "red",
    "blue",
    "green",
    "orange",
    "purple",
    "brown",
    "pink",
    "gray",
    "cyan",
    "magenta"
]

```

```

    "blue",
    "green",
    "orange",
    "purple",
    "brown",
    "pink",
    "gray",
    "olive",
    "cyan",
]
for i in range(optimal_k):
    mask = cluster_labels == i
    axes[1, 0].scatter(
        X_pca[mask, 0],
        X_pca[mask, 1],
        c=colors[i],
        label=f"Cluster {i}",
        alpha=0.6,
        s=50,
    )

# Plot cluster centers in PCA space
centers_pca = pca.transform(kmeans_optimal.cluster_centers_)
axes[1, 0].scatter(
    centers_pca[:, 0],
    centers_pca[:, 1],
    c="black",
    marker="x",
    s=200,
    linewidths=3,
    label="Centroids",
)
axes[1, 0].set_xlabel("First Principal Component (PC1)", fontsize=12)
axes[1, 0].set_ylabel("Second Principal Component (PC2)", fontsize=12)
axes[1, 0].set_title(
    f"K-Means Clustering (k={optimal_k}) in PCA Space", fontsize=14, fontweight="bold"
)
axes[1, 0].legend()
axes[1, 0].grid(True, alpha=0.3)

# Plot 4: Actual Target Labels in PCA Space
for i in [0, 1]:
    mask = y == i
    label = "No Heart Disease" if i == 0 else "Heart Disease"
    axes[1, 1].scatter(
        X_pca[mask, 0], X_pca[mask, 1], c=colors[i], label=label, alpha=0.6, s=50
    )
axes[1, 1].set_xlabel("First Principal Component (PC1)", fontsize=12)
axes[1, 1].set_ylabel("Second Principal Component (PC2)", fontsize=12)
axes[1, 1].set_title("Actual Labels in PCA Space", fontsize=14, fontweight="bold")
axes[1, 1].legend()
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("../assets/heart_disease_kmeans_clustering.png", dpi=150, bbox_inches="tight")
print("Visualization saved as '../assets/heart_disease_kmeans_clustering.png'")
plt.show()


```

=====

#### STEP 4: Visualizing Clusters

=====

Visualization saved as '../assets/heart\_disease\_kmeans\_clustering.png'



## Additional Analysis

Compare the discovered clusters with actual labels and analyze PCA components.

```
In [8]: print("\n" + "=" * 60)
print("Additional Analysis")
print("=" * 60)

# Compare clusters with actual labels
print("\nCluster vs Actual Label Cross-tabulation:")
crosstab = pd.crosstab(cluster_labels, y, rownames=["Cluster"], colnames=["Actual"])
print(crosstab)

# Feature importance in PCA
print("\nPCA Component Analysis:")
feature_names = X.columns
components_df = pd.DataFrame(
    pca.components_.T, columns=["PC1", "PC2"], index=feature_names
)
print(components_df.round(4))

print("\n" + "=" * 60)
print("Experiment Complete!")
print("=" * 60)
```

```
=====
Additional Analysis
=====
```

Cluster vs Actual Label Cross-tabulation:

Actual	0	1
Cluster		
0	147	37
1	13	100

PCA Component Analysis:

	PC1	PC2
age	0.2859	-0.4187
sex	0.1168	0.4316
cp	0.2862	0.1525
trestbps	0.1678	-0.3915
chol	0.0835	-0.4282
fbs	0.0761	-0.2399
restecg	0.1459	-0.2665
thalach	-0.3927	-0.0541
exang	0.3331	0.2083
oldpeak	0.3970	0.0617
slope	0.3520	0.0745
ca	0.3064	-0.1578
thal	0.3462	0.2634

```
=====
Experiment Complete!
=====
```